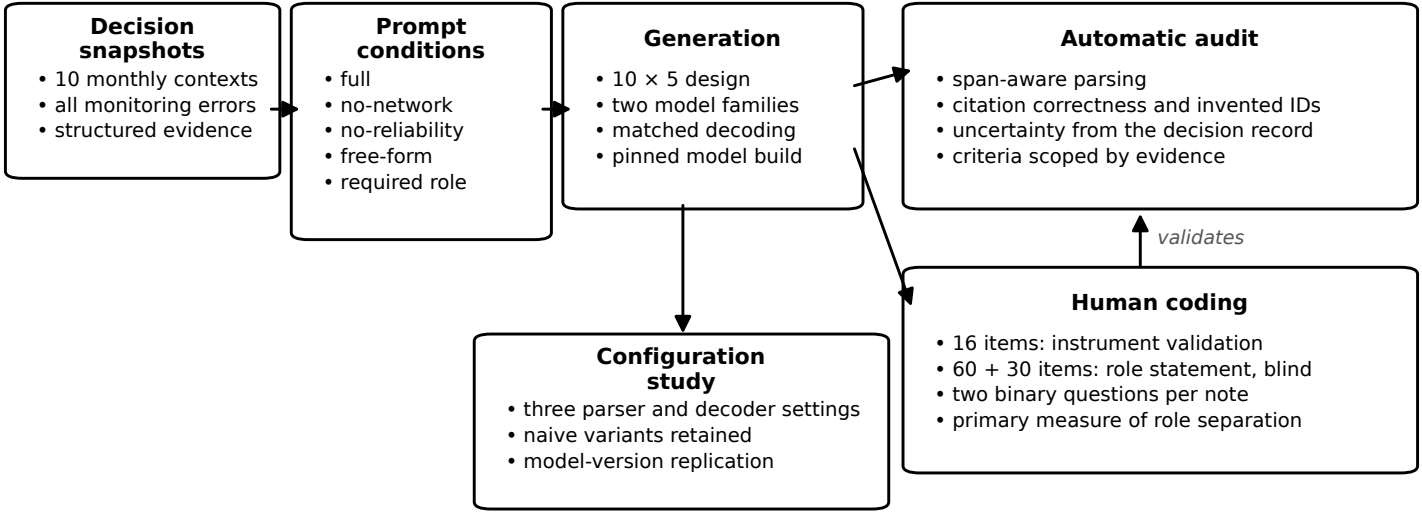
Decision snapshots
• 10 monthly contexts
• all monitoring errors
• structured evidence
Prompt conditions
• full
• no-network
• no-reliability
• free-form
• required role
Generation
• 10 × 5 design
• two model families
• matched decoding
• pinned model build
Automatic audit
• span-aware parsing
• citation correctness and invented IDs
• uncertainty from the decision record
• criteria scoped by evidence
validates
Human coding
• 16 items: instrument validation
• 60 + 30 items: role statement, blind
• two binary questions per note
• primary measure of role separation
Configuration study
• three parser and decoder settings
• naive variants retained
• model-version replication